


Centrality resolution vs true b (lower is better)



DeepSets/MLP-pool: lab-filtered subset (no track in $0 < \eta_{\text{lab}} < 1.5$ removed); MLP/Glauber/Truth: full set. Each ML model on its own test split. Not apples-to-apples — see script docstring.